

High-Level Programming Models for Heterogeneous Parallel Computing

Wen-mei Hwu

University of Illinois, Urbana-Champaign



Agenda

- Many-core Usage Pattern and Workflow
- Quick overview of current tools/environments
- Raising the Level of Kernel Development
- Conclusion and Outlook

GPU computing is catching on.

Financial
Analysis

Scientific
Simulation

Engineering
Simulation

Data
Intensive
Analytics

Medical
Imaging

Digital
Audio
Processing

Digital
Video
Processing

Computer
Vision

Biomedical
Informatics

Electronic
Design
Automation

Statistical
Modeling

Ray
Tracing
Rendering

Interactive
Physics

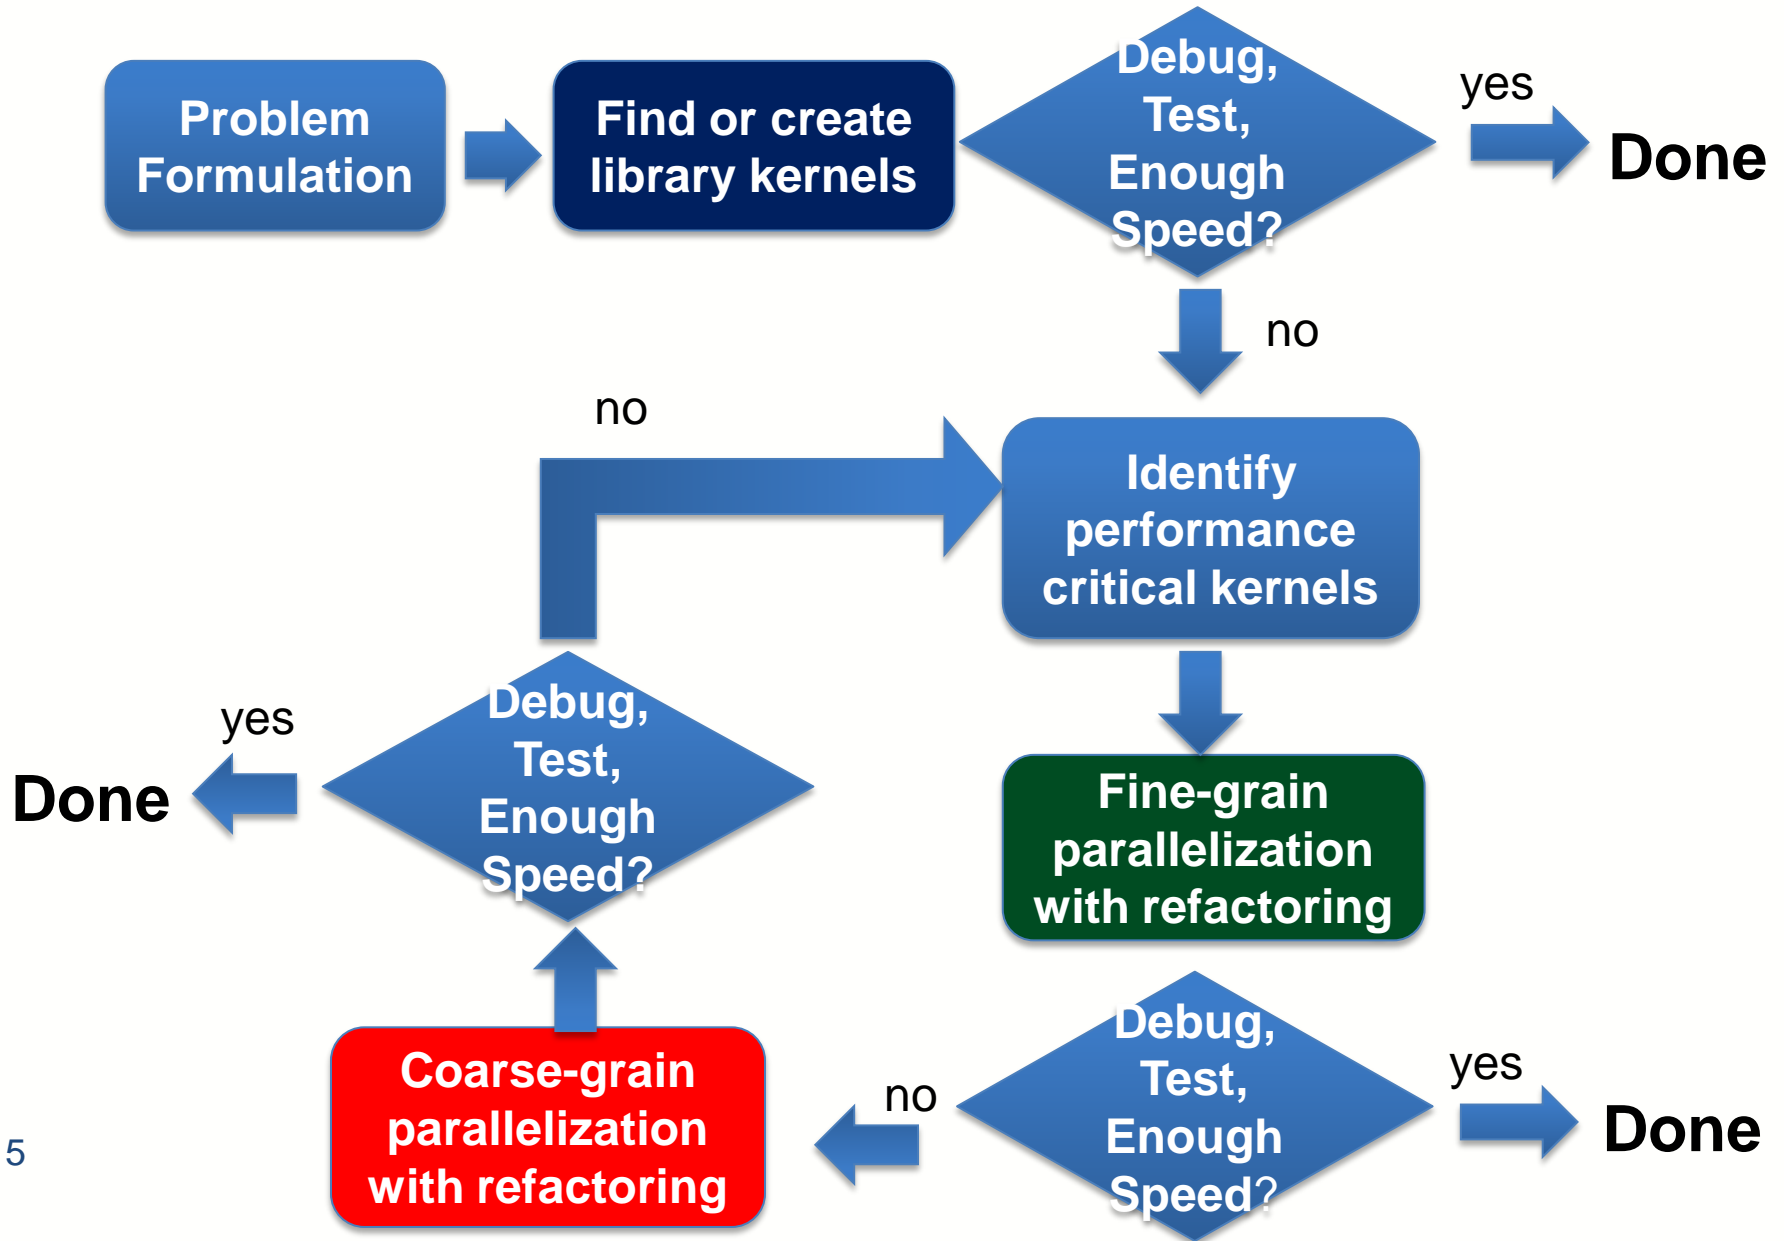
Numerical
Methods

- 280 submissions to GPU Computing Gems
 - 110 articles included in two volumes

A Common GPU Usage Pattern

- A desirable approach considered impractical
 - Due to excessive computational requirement
 - But demonstrated to achieve domain benefit
 - Convolution filtering (e.g. bilateral Gaussian filters), De Novo gene assembly, etc.
- Use GPUs to accelerate the most time-consuming aspects of the approach
 - Kernels in CUDA or OpenCL
 - Refactor host code to better support kernels
- 4 • Rethink the domain problem

Practical Parallelization Flow



Need for Better Kernel Interface

- Tedious and error prone
 - Each object has two versions
 - Programmers try to avoid copying by tracking which object is needed by CPU and GPU at each program point
 - Double buffering and triple buffering in copying used to hide latency
- In real applications
 - Disk I/O and MPI messages are often done by different threads than compute to hide latency
 - Data copying involves tricky thread synchronization

Using Kernels in Higher-Level Environments

- Thrust – C++ STL for CUDA
 - Containers, iterators, common kernels
 - <http://code.google.com/p/thrust/>
- Jacket- use of CUDA functions in MatLab code
 - Auto. GPU memory management and data transfer
 - <http://wiki.accelereyes.com>
- PyCUDA- use of CUDA functions in Python code
 - Consistent with numpy usage
 - <http://mathematician.de/software/pycuda>
- GMAC – Global shared address space
 - AutoGPU memory management and data transfer.
 - <http://adsm.googlecode.com/>

Where are the pain points?

- Currently, these systems do not use advanced compiler technology or hardware support
 - Implemented mostly through templates/macros
- Unnecessary data transfers before and after kernel execution
 - Could be optimized away with information on side-effect of kernels and live range of data objects
 - Portability for fusion architectures
- Double and triple buffering done by hand
 - Could be automated with compiler transformations

AVAILABLE KERNELS

Library Kernels

- CUBLAS
 - Basic Linear Algebra
 - CUDA SDK
- CULA, Magma
 - Linear Algebra Solvers
 - www.culatools.com
 - <http://icl.cs.utk.edu/magma>
- CUSP
 - Sparse data structures and algorithms
 - SpMV, CG, ...
- Graph algorithms
 - BFS kernels exist
 - Need graph partitioning kernels
- Unstructured grid algorithms
 - 3D surface mesh generation/refinement
 - Need 3D volume mesh generation (e.g. CGAL)/refinement
- Add your favorite library here

Four Challenges

- Computations with no known scalable parallel algorithms
 - Shortest path, Delaunay triangulation, ...
- Data distributions that cause catastrophic load imbalance in parallel algorithms
 - Free-form graphs, MRI spiral scan
- Computations that do not have data reuse
 - Matrix vector multiplication, ...
- Algorithm optimizations that are hard and labor intensive
 - Locality and regularization transformations

Kernel development for GPUs is heavy lifting.

Each kernel is typically a 3-month job but very few developers benefit from advanced compiler technology today.



12 Little code reuse due to kernel sensitivity to memory access patterns and work distribution.

KERNEL DEVELOPMENT

Many-core Kernel Development

- Many-core programming is about performance and scalability.
 - Scalability is also key to power efficiency.
 - Performance and scalability for many-cores requires largely the same techniques.
 - To regularize work and data for massively parallel execution.
 - To localize data for conserving memory bandwidth
- There is a gap between what the programmers need and what the tools provide today.

Key to Massive Parallelism - Regularity and Locality



Phillip Colella's "Seven dwarfs"

High-end simulation in the physical sciences = 7 numerical methods:

1. Structured Grids (including locally structured grids, e.g. Adaptive Mesh Refinement)
 2. Unstructured Grids
 3. Fast Fourier Transform
 4. Dense Linear Algebra
 5. Sparse Linear Algebra
 6. Particles
 7. Monte Carlo
- If add 4 for embedded, covers all 41 EEMBC benchmarks
 8. Search/Sort
 9. Filter
 10. Combinational logic
 11. Finite State Machine

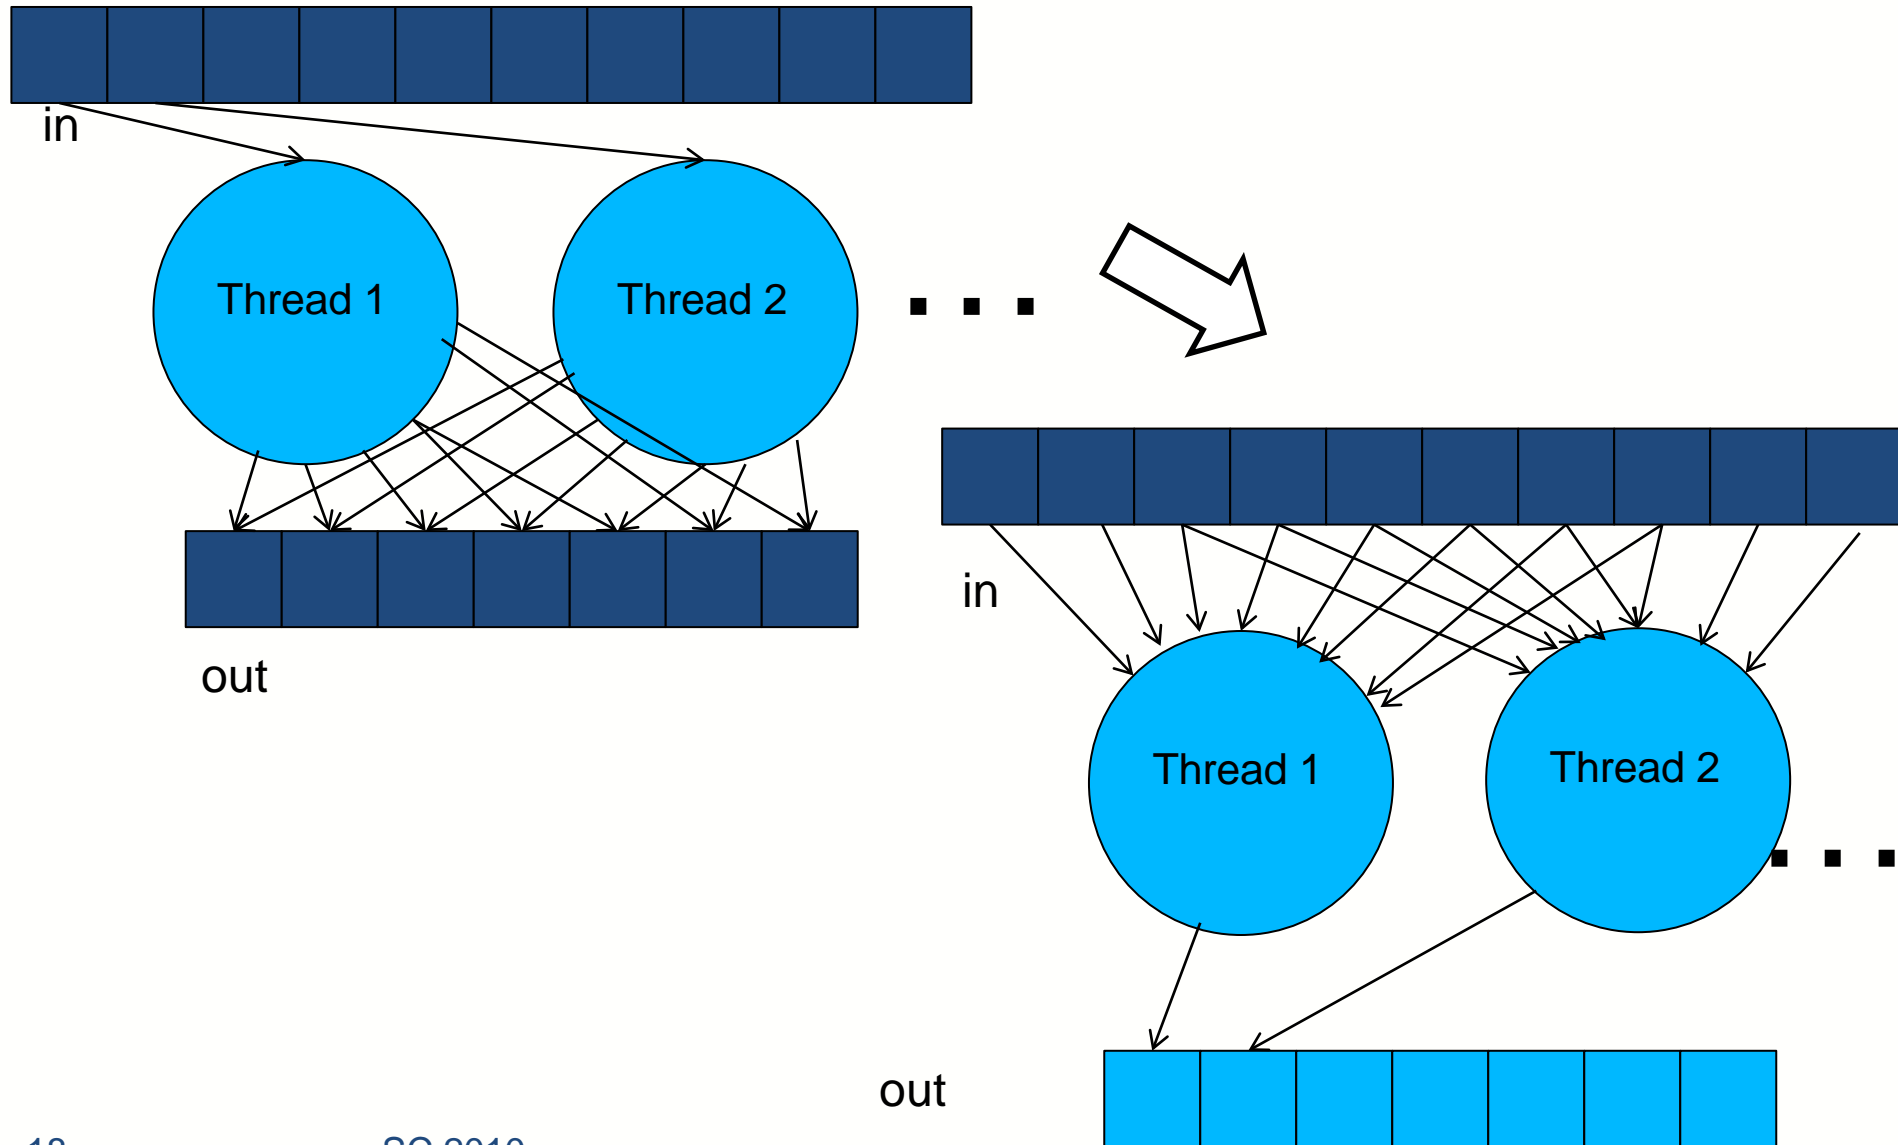
Well-defined targets from algorithmic, software, and architecture standpoint

Eight Optimization Patterns for Algorithms (so far)


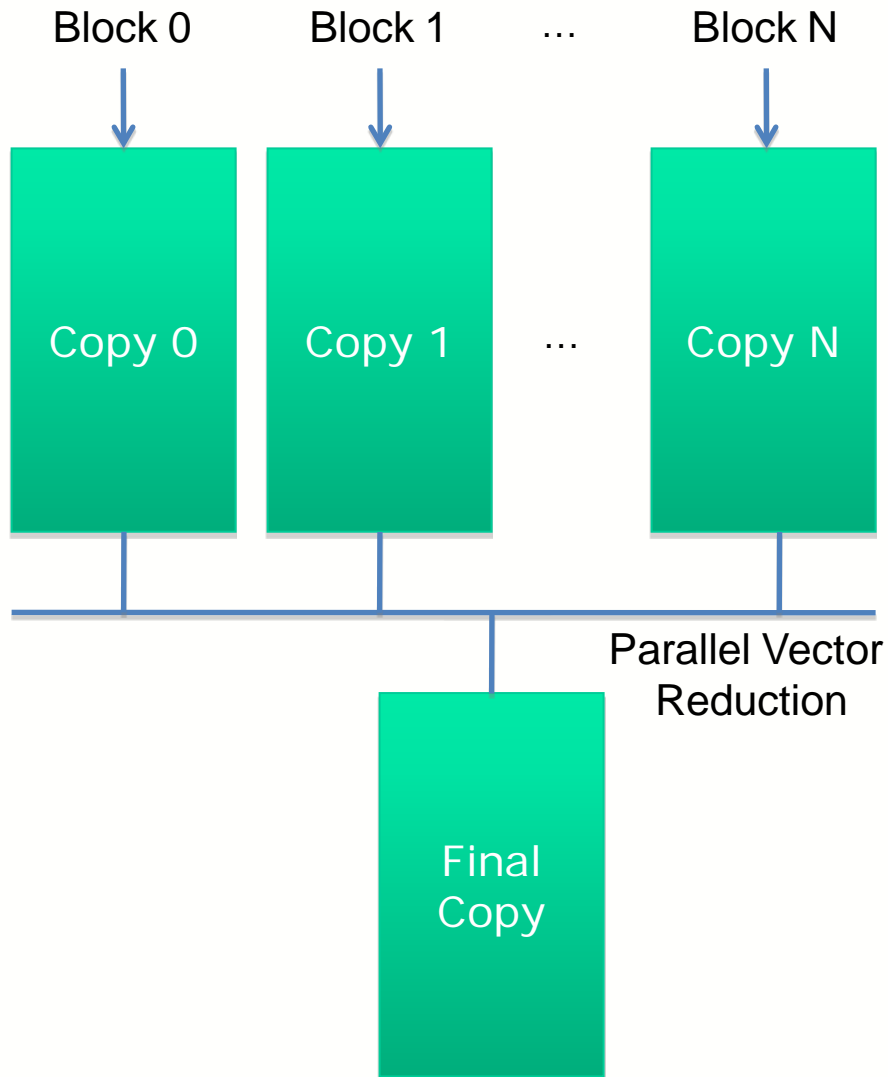
Technique	Contention	Bandwidth	Locality	Efficiency	Load Imbalance	CPU Leveraging
Tiling		X	X			
Privatization	X		X			
Regularization				X	X	X
Compaction		X				
Binning		X	X	X		X
Data Layout Transformation	X		X			
Thread Coarsening	X	X	X	X		
Scatter to Gather Conversion	X					

<http://courses.engr.illinois.edu/ece598/hk/>
GPU Computing Gems, Vol. 1 and 2

1: Scatter to Gather Transformation




2. Privatization



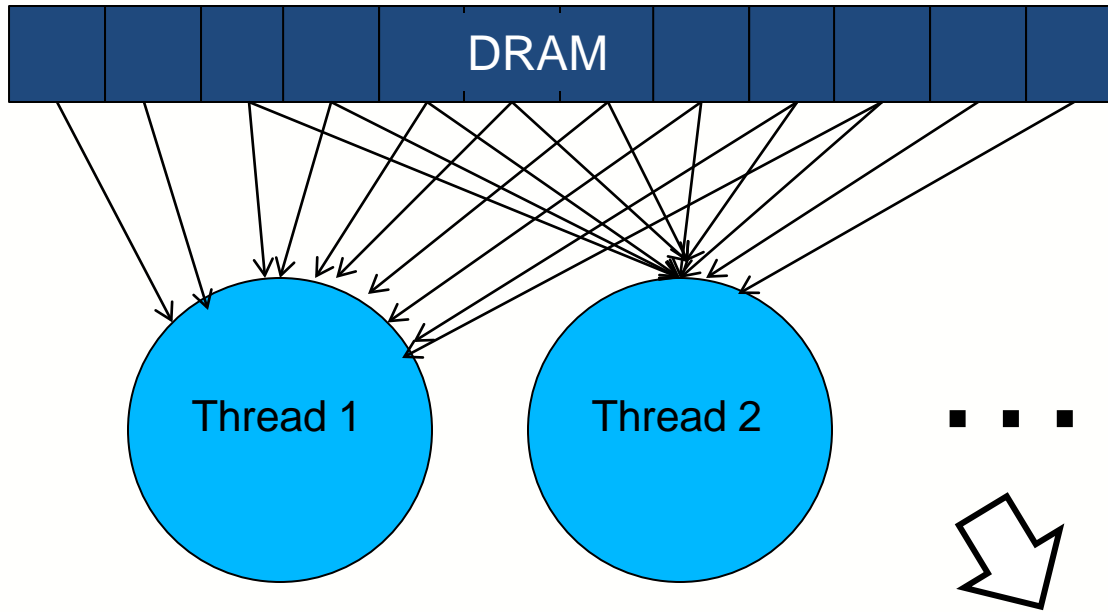
3. Granularity Coarsening and Register Data Reuse

- Parallel execution often requires redundant and coordination work
 - Merging multiple threads into one allows reuse of result, avoiding redundant work




4: Data Access Tiling


in




in



5. Data Layout Transformation




Array of Structure: $[z][y][x][e]$







Structure of Array: $[e][z][y][x]$

4X faster than AoS on GTX280




6: Input Data Binning




7. Compaction





Variable sized bins,
sort and scan




CPU



On-Chip Memory



8. Regularization



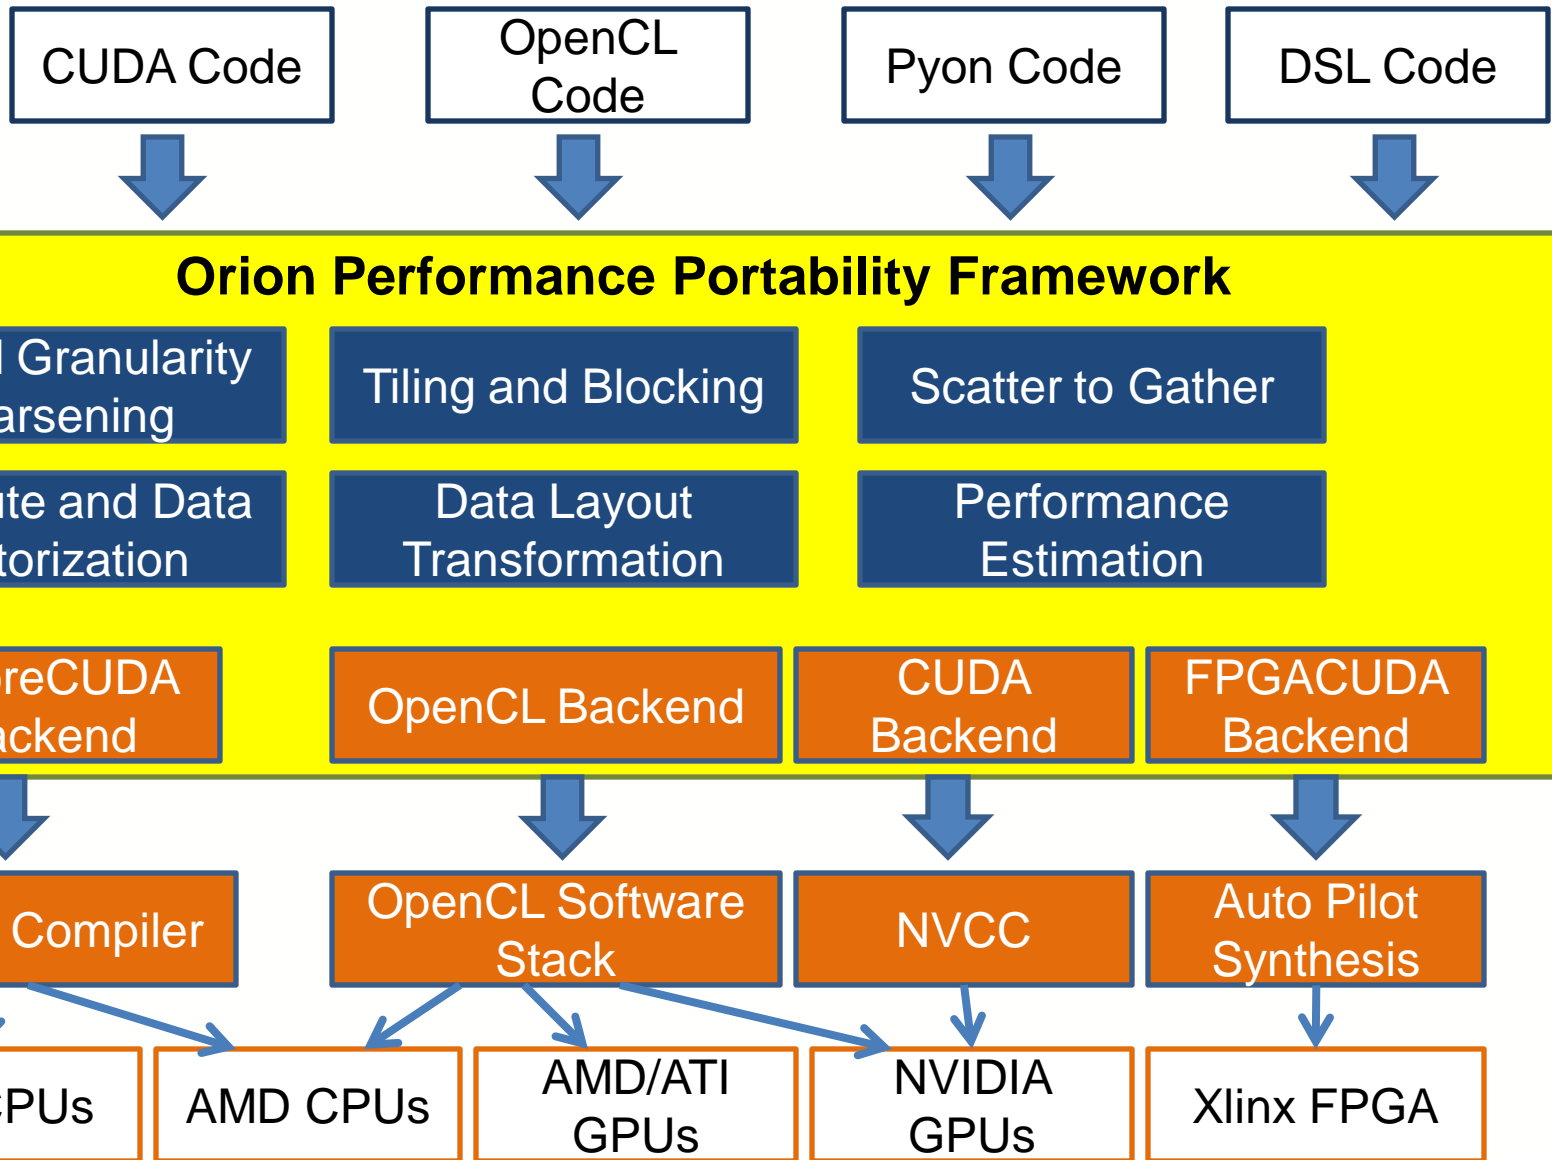
Tools go with techniques.

- Tools should facilitate key techniques
 - Programmers should write code “for others to understand instead of for computers to execute”
 - Dijkstra
- Techniques vary in their potential for automation
 - Scatter-to-gather, granularity coarsening, data access tiling, and memory layout quite amenable
 - Need clear performance guidance
 - Input binning, bin sorting, and hierarchical queues are much harder
 - Need to provide APIs understood by compilers/tools
 - Developer feedback critical to success

Some Current Efforts

- PGI FORTRAN (PGI)
 - FORTRAN loop compilation into kernels
- Chapel (CRAY)
 - Chapel vector loop compilation into kernels
- Copperhead (Berkeley)
 - Annotated Python subset JITed into kernels
 - Modest performance within interpreter use model
- Pyon (Illinois)
 - Strongly typed Python subset compilation into kernels
 - Aggressive memory transformations


Orion: Reducing Performance Cost of Heterogeneous Parallelism



Conclusion and Outlook

- Standard and domain library kernels are main use model of many-cores.
 - Kernel development is heavily lifting
 - Useful abstractions will require compilers to do some heavy lifting.
- However, compilers/tools are fragile.
 - Compilers transformations need to be part of the development, rather than afterthought
 - Developers must be able to reach into the abstraction whenever they want.

Crossing the Valley of Death



By giving developers what they want.



Acknowledgements

- D. August (Princeton), S. Baghsorkhi (Illinois), N. Bell (NVIDIA), D. Callahan (Microsoft), J. Cohen (NVIDIA), B. Dally (Stanford), J. Demmel (Berkeley), P. Dubey (Intel), M. Frank (Intel), M. Garland (NVIDIA), M. Gschwind (IBM), R. Hank (Google), J. Hennessy (Stanford), P. Hanrahan (Stanford), M. Houston (AMD), T. Huang (Illinois), D. Kaeli (NEU), K. Keutzer (Berkeley), I. Gelago (UPC), B. Gropp (Illinois), D. Kirk (NVIDIA), D. Kuck (Intel), S. Mahlke (Michigan), T. Mattson (Intel), N. Navarro (UPC), J. Owens (Davis), D. Padua (Illinois), S. Patel (Illinois), Y. Patt (Texas), D. Patterson (Berkeley), C. Rodrigues, S. Ryoo (ZeroSoft), C. Rodrigues (Illinois) K. Schulten (Illinois), B. Smith (Microsoft), M. Snir (Illinois), I. Sung (Illinois), P. Stenstrom (Chalmers), J. Stone (Illinois), S. Stone (Harvard) J. Stratton (Illinois), M. Valero (UPC)
- And many others!

THANK YOU!